

# Review of RISC-V (and other) zkVMs



by Eduard Sanou from PSE (EF)



# Overview

- RISCv intro
- What is a zkVM
  - A note on privacy
- Comparison Overview
  - Types of zkVM
- Architecture
  - Let's take a look
  - Continuations
  - Coprocessor
- Benchmarks



# RISC-V intro

- Open CPU architecture following the Reduced Instruction Set Computer paradigm
  - Simpler than ARM64
- 32 bit and 64 bit flavors
- Multiple standard extensions
  - I: Base Integer ISA
  - M: Integer Multiplication & Division
  - A: Atomic
  - C: Compressed
  - F: Floating Point
- Support from major compilers: gcc LLVM (Rust, clang, ...), go lang

```
.section .text
.global strlen
strlen:
    # a0 = const char *str
    li    t0, 0          # i = 0
1: # Start of for loop
    add   t1, t0, a0      # Add the byte offset for str[i]
    lb    t1, 0(t1)       # Dereference str[i]
    beqz   t1, 1f          # if str[i] == 0, break for loop
    addi   t0, t0, 1       # Add 1 to our iterator
    j      1b             # Jump back to condition (1 backwards)
1: # End of for loop
    mv     a0, t0          # Move t0 into a0 to return
    ret                    # Return back via the return address register
```



# What is a zkVM

- VM: Machine that behaves like a computer. Can do arithmetic ops, bitwise ops, branching, read/write memory, etc.
- zkVM: a system that can generate zk proofs of correct execution under a VM.
  - Execution trace: chronological list of instructions (with input-output) that are executed when running a VM program.
  - A zkVM requires: a circuit that verifies the correctness of an execution trace of the VM.
    - Realized by repeating the following verifying logic on each trace step:
      - Fetch, Decode?, Execute, Memory I/O



# A note on privacy

- Is there really zk in these zkVMs?
- No privacy in most of the zkVM projects, so the term zk can be misleading: **there's no zero knowledge**.
- Usually the goal is **succintness** (the "S" from SNARK).

“Verifiable computing (colloquially mis-termed ZK) is an incredibly powerful technology” [1]



# Comparison Overview

Project	ISA	Arithmetization	circuit framework	Proof System	Field
RiscZero	RV32IM	AIR+Plonkish	zirgen [10]	STARK	BabyBear
SP1	RV32IM	AIR	Plonky3	STARK	BabyBear
Powdr	RV32IMAC [1]	Plonkish [5]	custom	pluggable [6]	pluggable [7]
Jolt	RV32I	R1CS+lookups	custom	Spartan + Lasso	BN254 scalar
zkm	MIPS32	AIR	Plonky2	STARK	Goldilocks
o1VM	MIPS32	Plonkish	Kimchi	Kimchi (PLONK)	BN254 scalar
zkWASM	WASM	Plonkish	halo2	halo2	BN254 scalar
Nexus zkVM	NexusVM [3][4]	R1CS	custom	(Hyper)Nova	255 bits [9]
CairoVM	Cairo	AIR	custom	STARK	felt252
Valida	Valida [2]	AIR	Plonky3	STARK	31 bits [8]
Triton VM	Triton [11]	AIR+AET	custom	STARK	Oxfoi [12]

[1] [https://github.com/powdr-labs/powdr/blob/3f771638467a8bd527805239ac614dd56e65f11c/riscv/tests/instruction\\_tests/README.md](https://github.com/powdr-labs/powdr/blob/3f771638467a8bd527805239ac614dd56e65f11c/riscv/tests/instruction_tests/README.md)

[2] <https://github.com/valida-xyz/valida/blob/898d5ce0581095eabdb5c7219f222cb17cdd46a9/opcodes/src/lib.rs>

[3] <https://docs.nexus.xyz/Specs/nexus-vm#nexus-virtual-machine-instruction-set>

[4] <https://github.com/nexus-xyz/nexus-zkvm/blob/a3a2b156b90276fc456a58f15fd46b638d2ca010/vm/src/instructions.rs#L22>

[5] Powdr uses a PIL-like language, which allows extensions to Air making it Plonkish

[6] Currently implemented: halo2, eSTARK. Work in progress: SuperNova.

[7] Goldilocks (63 bits) with eSTARK backend, 255+ bit fields with halo2 backend (for example: BN254 scalar field).

[8] Both BabyBear and Mersenne31 are 31 bit fields.

[9] Uses the BN254 and Grumpkin curve cycle

[10] zirgen is closed source. This tool allows compilation of circuits from a high level description; the output generated for RiscZero is public.

[11] <https://triton-vm.org/spec/instructions.html>

[12] Oxfoi is a 64 bit prime <https://triton-vm.org/spec/isa.html#oxfoi>



# Types of zkVM

Vitalik defined a classification for zkEVMs [1]. This is my proposal for zkVMs.

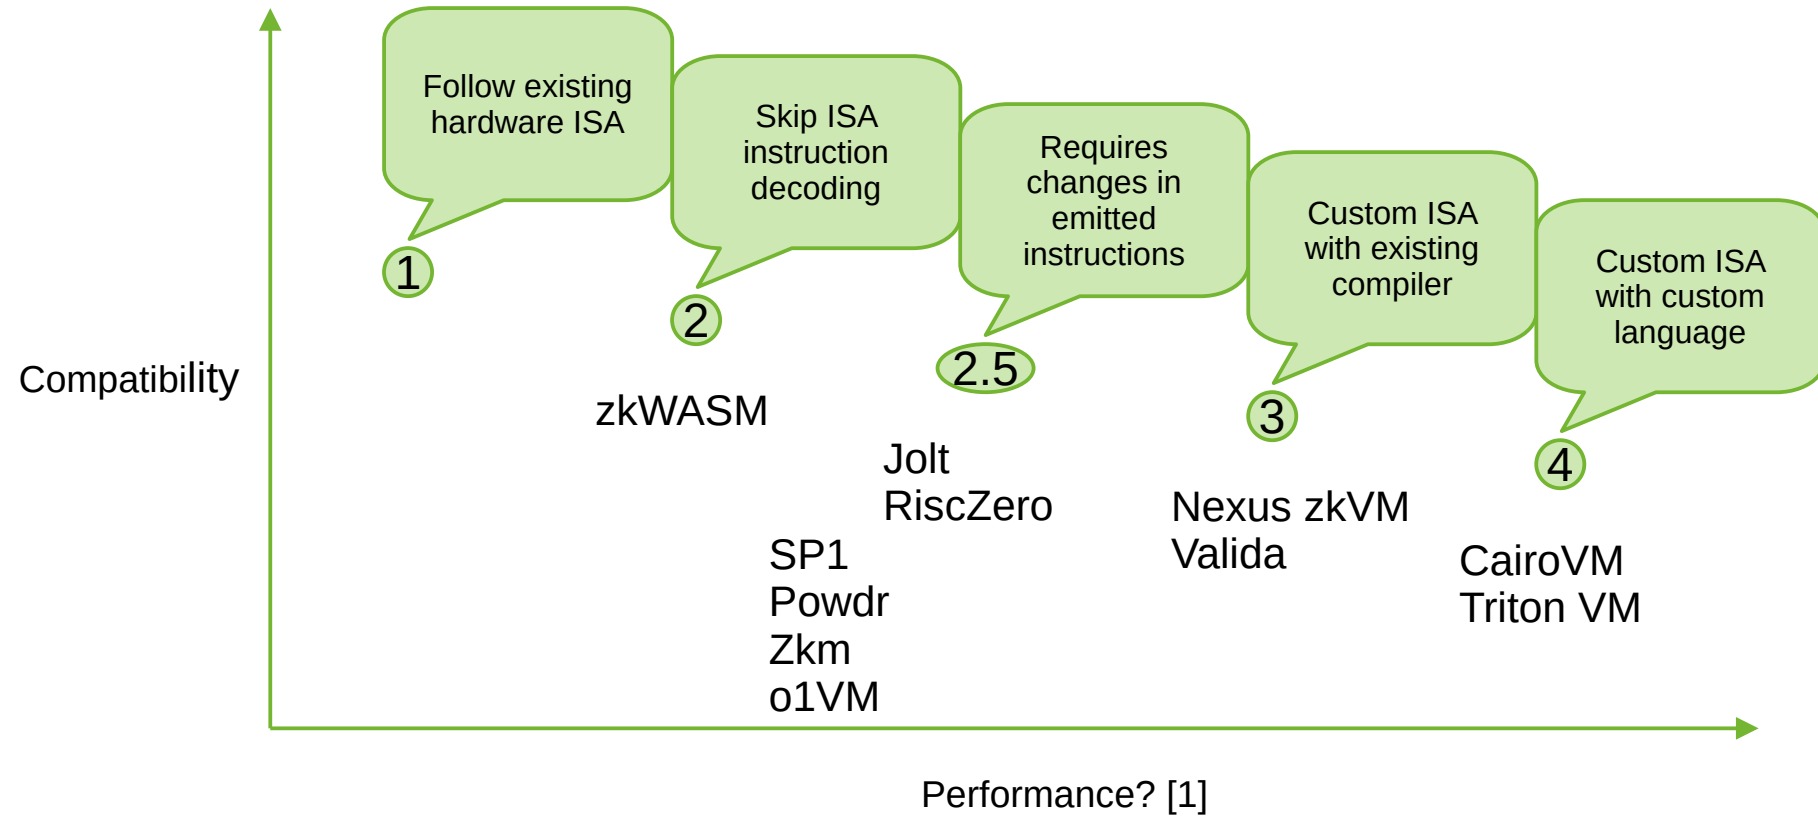
- **Type 1:** Follows the full spec of an existing hardware ISA (RISCV, MIPS, etc).
  - Includes instruction decoding after fetching from memory.
- **Type 2:** Encodes instructions in a custom format.
  - Decoding to custom format done during compilation
  - Can't write to memory and execute it (no JIT / AOT compilation)
- **Type 2.5:** Requires some changes on emitted instructions
  - Either via transformation in the decoding step or by patching the compiler.
- **Type 3:** Targets custom ISA (potentially zk friendly) but can be connected to existing compiler framework
  - Example: implements an LLVM backend for the custom ISA
- **Type 4:** Targets custom zk friendly ISA and requires custom high level language

[1] <https://vitalik.eth.limo/general/2022/08/04/zkevm.html>

Note 1: Vitalik's categories distinguish matching gas cost. We could translate that to CPU cycles, but the RISCV ISA doesn't specify the cycles per instruction (that's implementation dependent).



# Types of zkVM



[1] Justin Thaler argues that a zk friendly ISA doesn't bring better performance  
<https://a16zcrypto.com/posts/article/faqs-on-jolts-initial-implementation/#section--8>





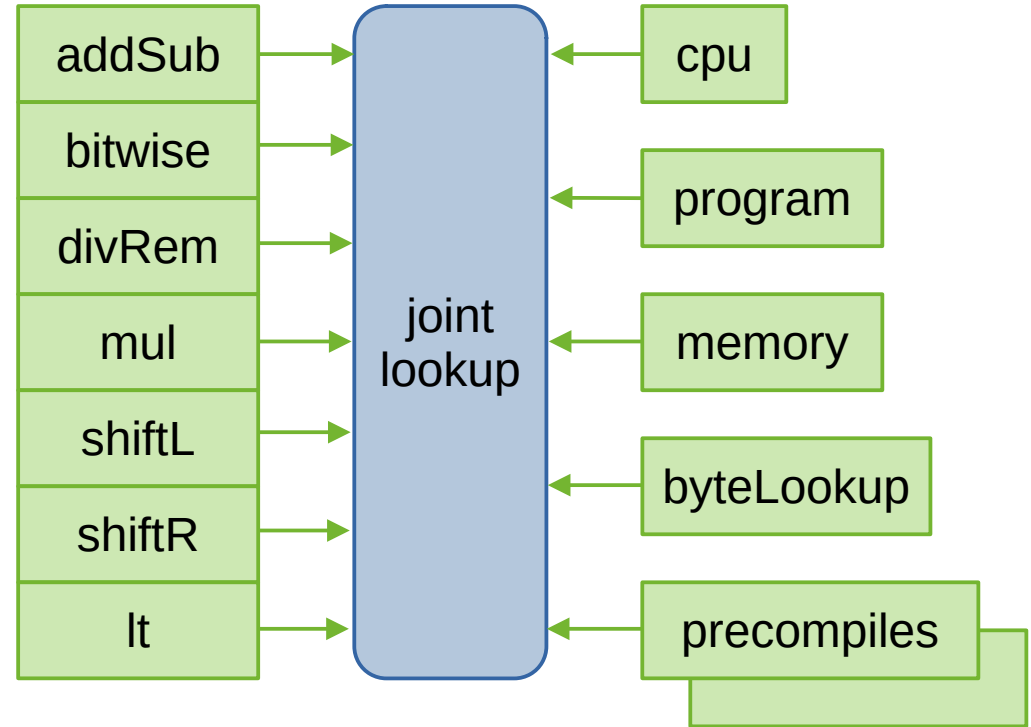
# Architecture

- Main blocks:
  - CPU: “Loop” logic that performs Fetch, Decode?, Execute, Memory I/O
    - Execute sometimes delegated to other blocks
  - Memory: Verify RAM read/write consistency
- Extras:
  - “std”: heap, stdin, stdout, threads, time, file system, ...
  - Continuations
  - Coprocessors



# Architecture example: SP1

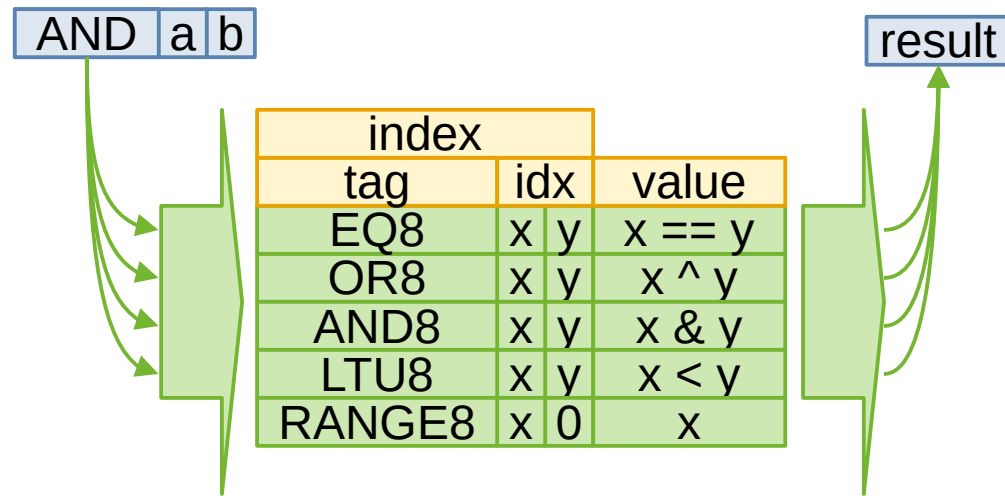
- Each chip defined with AIR constraints
- Chips do IO via joint lookup argument
- Cpu [1] is the entry point. AIR constraints verify 1 step per row:
  - Fetch instruction from program
  - Prepare instruction arguments
    - Maybe read from memory
  - Apply instruction
    - Possibly via lookup to another chip
  - Update Cpu state (registers, pc, timestamps, clk)
  - Maybe write result to memory



[1] <https://github.com/succinctlabs/sp1/blob/5db203c55647c30618431822d33f419614f9fab6/core/src/cpu/air/mod.rs#L32>

# Architecture example: Jolt [1]

- Built with Lasso lookups
- Each RISC-V instruction is defined with a decomposable lookup:
  - Decompose arguments into bit chunks
  - Use chunks for indexed lookups to small subtables
    - Tables need to be MLE [3]
  - Combine lookup results
- Merge all lookups, leading to a single decomposable lookup per instruction.
  - Merge all subtables into one
  - Merge all combine expressions into one
- Use R1CS for glue and some instructions



[1] <https://people.cs.georgetown.edu/jthaler/Jolt-paper.pdf>

[2] <https://people.cs.georgetown.edu/jthaler/Lasso-paper.pdf>

[3] Table must be the evaluation of a Multi Linear Expression: an expression that takes n bits and can be evaluated in  $O(n)$

# Continuations

- Long computation may not be feasible in a single proof
  - High memory requirements
  - Proof system limitations
- Split the computation trace into chunks and generate a proof for each one
  - Possibly aggregating the proofs later for constant final proof size
  - Requires connecting the execution context between continuations
- Supported by: RiscZero, zkm, zkWASM, SP1, Powdr, zkWASM, Triton VM, Nexus (via folding) [1]

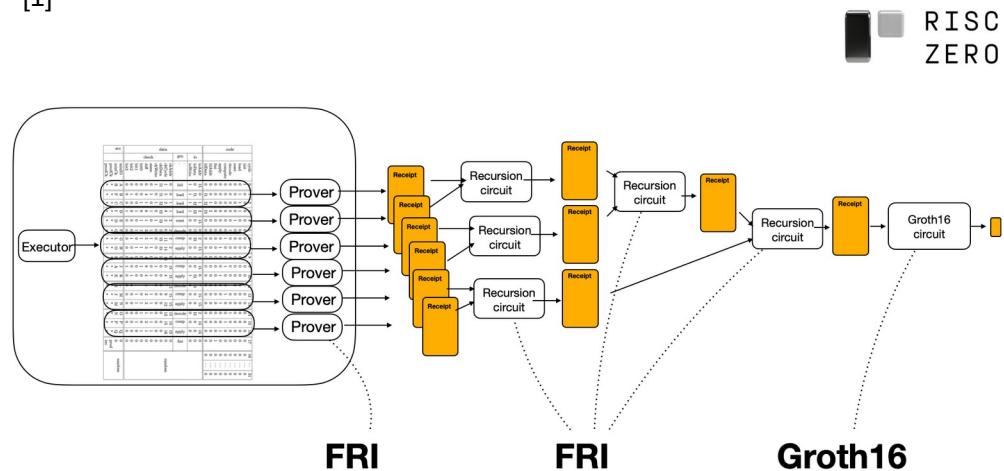
[1] Probably all zkVMs will end up supporting continuations



# Continuations example: RiscZero

- Split execution trace into segments of equal number of cycles (steps)
- Between segments, memory is passed via a Merkle Tree in pages
- Inside segment:
  - Beginnig: used pages are “decommitted” from MT
  - During: memory ops are verified via permutation check + sorted table [2]
  - End: used pages are “committed” to MT
- Each segment is proved independently (STARK w/ RISC-V Circuit)
- All segment proofs are aggregated recursively (STARK w/ Recursion Circuit)
- Recursion proof can be verified via Groth16 for on-chain verification (Groth16 w/ Groth16 Circuit)

[1]



[1] <https://dev.risczero.com/proof-system/>

[2] “sorting paradigm” in <https://eprint.iacr.org/2023/1555.pdf>



# Coprocessor

Also called “precompiles”

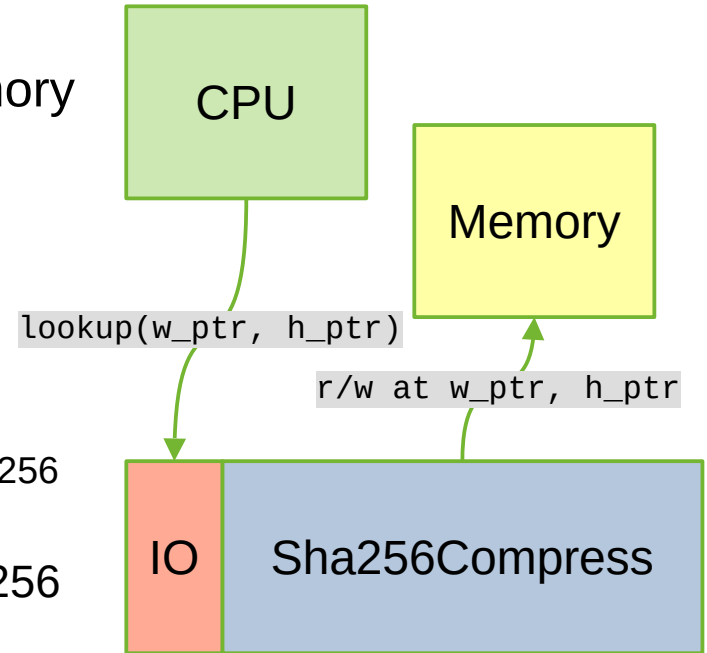
- Some programs can take a lot of steps when broken down to RISC-V instructions
  - Can we accelerate the bottlenecks with custom circuits?
- Examples:
  - Cryptographic operations that require emulated fields (EC signatures, pairings, etc.)
  - Hashing (Keccak, Sha256, Poseidon, etc.)
- Supported by [1]: SP1, Nexus, Powdr,

[1] Probably all zkVMs will end up supporting coprocessors



# Coprocessor example: SP1 [1]

- ISA level: accessed via system calls
- Circuit level: implemented with new AIR that shares memory with main VM. Accessed via lookup [2] with IO table.
- Example for sha256 compress [3]:
  - Cpu chip:
    - Lookup with *input word memory pointer*, *sha256 state memory pointer*.
    - Bumps memory timestamp by the number of memory ops the sha256 compress will do.
  - Sha256Compress chip: receives lookup and performs sha256 compression:
    - absorb word (read at *input word memory pointer*)
    - do hash operations
    - update state (read + write at *sha256 state memory pointer*)



[1] <https://succinctlabs.github.io/sp1/writing-programs/precompiles.html>

[2] Technically it's a permutation check

[3] <https://github.com/succinctlabs/sp1/blob/5db203c55647c30618431822d33f419614f9fab6/core/src/syscall/precompiles/sha256/compress/air.rs#L27>

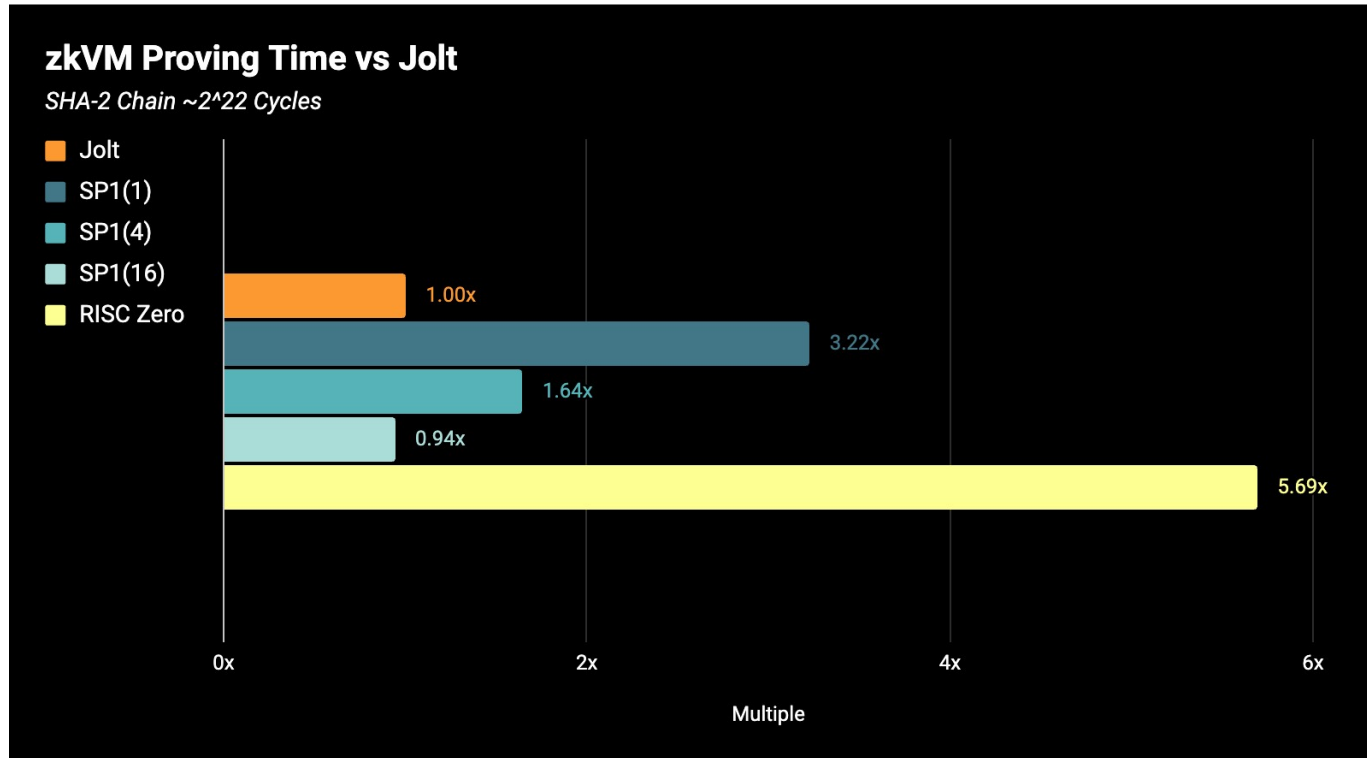
# Benchmarks

- WARNING: Benchmarking is tricky because there are many subtleties:
  - CPU vs GPU implementations
  - Better/worse parallelism
  - Some zkVMs support continuations and other don't
  - Some zkVMs support recursion and others don't
  - Each zkVM may be in a different stage of optimization





# (Partial) Benchmarks



Note: Value in parenthesis is number of shards



# Benchmarks

- Jolt can prove 100k instr/sec.
  - “[...] over twice the on-board computing power of the Apollo 11 mission.”
  - About 30x slower than a TI-84 graphing calculator.

Quote and numbers from Jolt authors at <https://a16zcrypto.com/posts/article/building-jolt/> based on benchmarks run on an Apple M3 Max

[1] Jolt article says “150x slower”. But TI-84 uses a z80 @ 15Mhz; and in z80 each instruction takes 4+ cycles.





Questions?

